

① RECORD CORRECTION — TAO Halving Date

A prior archive PDF contained a factual error. It stated the second halving would occur *in late 2026 or 2027* — off by 2-3 years. Source confirmed via Taostats.io (official block explorer).

Halvening	Date	Block Reward	TAO Supply
H1 — First	Dec 15, 2025	0.5 TAO	10,500,000
H2 — Second ♦	Dec 12, 2029	0.25 TAO	15,750,000
H3	Dec 10, 2033	0.125 TAO	18,375,000
H4	Dec 7, 2037	0.0625 TAO	19,687,500

Halvings are supply-based (every 10.5M TAO issued) — approximately every 4 years. **This record supersedes all prior references.**
Next halving: **December 12, 2029.**

② MANTIS (SN123) — Decentralized Signal Refinery

Source: TAO Daily — How MANTIS Orchestrates a Coordinated Pipeline for Intelligent Trade Execution. **Relevance to TaoBot:** HIGH.

- SN123 is a decentralized forecasting subnet — miners submit prediction embeddings; validators score by **marginal information gain**. No marginal gain = no reward. It prices signal quality, not opinion.
- 4-layer pipeline: Upstream subnets (SN13, 33, 6, 22, 50, 82) → MANTIS filters signal → Meta-models (direction, regime, volatility) → **Vanta (SN8)** executes under risk constraints.
- Architecture is *cyclical*: some subnets feed upstream priors into MANTIS and later consume MANTIS outputs downstream — reinforcing feedback, not a linear pipe.
- **Vanta (SN8)** is the execution endpoint — already doing risk-gated trade selection. Future TaoBot integration candidate.
 - *TaoBot's internal signal layer should weight strategies by marginal gain, not equally. A signal that doesn't improve the ensemble prediction gets deprioritized — no manual tuning required.*

③ TEUTONIC (SN3) — Bittensor's Answer to Covenant

Source: TAO Daily — Teutonic (SN3) Is Cooking a 24B Looped Transformer. **Relevance to TaoBot:** MEDIUM-HIGH.

- Const rebuilt SN3 from scratch **4 days after Covenant AI abandoned Templar**. Renamed Teutonic. Seed king: 0.9B Gemma3, launched April 13, 2026. Loss: ~13 → low 5s through open competition.
- **24B Looped Transformer:** Reuses the same weight block multiple times per forward pass. Reasoning depth becomes an inference-time knob — more loops = deeper reasoning without more training parameters.
- ByteDance's Ouro: 1.4B looped model performing like 12B on benchmarks. Claude Mythos suspected to use same architecture (~80% on GraphWalks BFS vs GPT-5's 21%).
- King-of-hill mechanism: hardware-agnostic — only cross-entropy loss matters. Opens the field to data curation, training tricks, and architectural innovation.

■ **DO NOT BUY SN3 alpha token — owner key situation unresolved. Const's own warning. TaoBot watchlist: SN3 stays YELLOW.**

■ *Looped architectures = inference-time compute scaling. When TaoBot adds AI-based signal generation, prefer models that reason deeper over models that are merely larger.*

④ THE ORCHESTRATOR PDF — Key Takeaways

Source: The Orchestrator.pdf — Authored by II Agent, April 15, 2026. 14 pages. Already in The Archives.

- **Architect vs Orchestrator:** The Architect designs once (12 strategies, OpenClaw, gate pipeline). The Orchestrator runs continuously — reads market, scores fleet, classifies regime, generates directives. Two distinct roles. Never confuse them.
- Everything described in the document is already live: II Agent 5-min cycle, BULL/BEAR/SIDEWAYS/VOLATILE regime detection, HOT/HEALTHY/WATCHING/STRUGGLING fleet tiers, 40-entry observation ring buffer, directive generator, OpenClaw BFT at 33.9% approval.
- **Section VIII — What's NOT yet built** (the real roadmap): Regime-aware position sizing · Consecutive loss circuit breaker · Performance trending detection · Cross-strategy correlation detection · OpenClaw-weighted regime · Validator performance weighting by subnet.
 - *Circuit breaker first: WR averages hide streaks. A 60% strategy can go 0-for-8 on a bad night. The consecutive loss circuit breaker catches what averages miss.*
 - *Cross-strategy correlation detection is subtle but critical — 10 bots agreeing BUY off the same MACD signal is one signal counted 10 times, not true consensus.*

⑤ HOSTING DECISION — PENDING

Platform	Always-On	RAM	No Card	Verdict
Render Free	✗ Sleeps	512MB	✓	Wrong for bot
Fly.io Free	✓	256MB	✗ Card	RAM too low
Oracle Always	✓	1GB	✗ Verify	Best free option
Vultr (\$6/mo)	✓	1GB+	✓ Crypto	Best paid option
Railway Hobby	✓	512MB	✗ Card	Easiest upgrade

Decision: Wife's credit card available. Railway Hobby (\$5/mo) is the fastest path. Vultr (\$6/mo, Bitcoin accepted) is the decentralized path. Migration guide is written — 1–2 hours to execute. **Do this at the start of the next coding session.**

“Architecture tells you how the pieces fit. Orchestration tells you how the pieces move. The Orchestrator is already moving. Every path leads to TaoBot.”